



Where the model looks: RISE maps on the deployed pipeline


RISE (Petsiuk et al. 2018, arXiv:1806.07421) covers a photograph with random coarse masks, scores every masked image with the frozen BioCLIP 2 tower and the linear probe, and averages the masks weighted by the probability of a target genus, so the map shows what the deployed black-box pipeline actually relies on. We ran it on our own pipeline to ask the question the accuracy numbers cannot answer: when the classifier names a genus, is it looking where a taxonomist would? Maps are qualitative - the blob scale is one cell of the 7×7 mask grid, and the weights use the probe's raw softmax because the deployed temperature-calibrated probabilities saturate at 1.0 and carry no signal - but they are how we check that a right answer is right for a defensible reason. Parameters: 400 masks, 7×7 grid, keep probability 0.5.




Odontomachus: the hotspot sits on the head capsule and the mandible bases, exactly the region a key uses to separate the trap-jaw genera.



Anochetus, the smaller trap-jaw look-alike: the same head-and-mandible region carries the decision on this specimen too.



Technomyrmex: the gaster is the hottest region - plausibly the gastral segmentation that separates it from *Tapinoma*, which on the same photo is read from the head instead.



Strumigenys, the probe's best genus, and an honest surprise: the map favours the head capsule and the gaster, not the famous trap jaws.